

HPBBM. Exercises R. Summaries, tables, aggregate, by.

Ramon Diaz-Uriarte*

2016-11-17 (Release: Rev: fe4c981)

Contents

A Recommendations	1
B Purpose of the exercises	1
C General hints	1
C.1 The data	2
C.2 Working directory	2
C.3 Grading	2
1 Tables and cross-tabs	3
2 Two subsetting operations	3
3 Gene summaries by condition and sex	4
3.1 aggregate: The median of three genes by condition	4
3.2 aggregate: The median of three genes by condition and sex	4
3.3 aggregate: The median of all the genes by condition and sex	4
3.4 aggregate: The mean and standard deviations of three genes by condition and sex	5
3.5 by and aggregate: those three genes again, but now use “summary”	5

A Recommendations

- Return the exercise as a single text file with commented code.
- Use comments if you think it is justified.

B Purpose of the exercises

The purpose of this exercise is to get used to obtaining summary statistics from data sets and from subpieces of those data sets.

C General hints

You will use the leukemia data (including the data, the class and the sex objects we used in the exercise for figures). Here are those instructions again:

*Dept. of Biochemistry, Universidad Autónoma de Madrid, Spain, <http://ligarto.org/rdiaz>, rdiaz02@gmail.com

C.1 The data

- Read the data “leukemia.data.txt”. Note that there are row names that we do not want to be part of the data itself (look at the help for “read.table”, the option “row.names”). Call this “leuk.dat”. These data are based on the famous leukemia data from Golub et al. I’ve modified a few things below: the gene names and the sex of the patients are invented.
- That will be a data frame. For most figures your life will be simpler if you convert it into a matrix, using `data.matrix`. Call this data matrix “leuk.dat.m”. (Why is it a good idea to convert this into a matrix? Try doing the figures using the original data frame.)
- Read the data “leukemia.class.txt”. Use `scan` for this (do you know why?). Since these are labels, use `what = ""` in `scan`. Convert the classes into a factor. Call this “leuk.class”.
- Create a vector for the sex of the patients. The patients are Male and Female, alternating, and starting with Male. Call this factor “sex”.

C.2 Working directory

All of your code should run. Create a directory called `R-ex-1`; place in this directory the data files and the script with code. I will run your scripts from this directory. So, when reading data, you do NOT have to specify paths. For instance, do not do

```
read.something("~/R-ex-1/somedata")
```

but do

```
read.something("somedata")
```

C.3 Grading

The grading is defined in the “Rúbrica”.

1 Tables and cross-tabs

Obtain a simple table of how many males and females are there. Use `table` on the “sex” object:

```
## sex
## Female    Male
##        19     19
```

Now, a table of how many of each sex are there in each class, or viceversa. This requires you to pass two arguments to `table`

```
##           leuk.class
## sex      ALL AML
## Female  13   6
## Male    14   5
```

(if your table looks transposed with respect to the above ... well, change the order of arguments).

Now, do the same as above, but with a formula interface, which some times is more useful. Use the `xtabs` function

```
##           leuk.class
## sex      ALL AML
## Female  13   6
## Male    14   5
```

Finally, rerun the code from `xtabs` and assign that to an object, call it “mytable” and convert mytable to a data frame using `as.data.frame` and print it:

```
as.data.frame(mytable)

##      sex leuk.class Freq
## 1 Female      ALL    13
## 2 Male      ALL    14
## 3 Female      AML     6
## 4 Male      AML     5
```

Notice that this is just a very simple but extremely powerful way of getting frequencies and saving them as a data frame for further analysis!! With the output from the tables we could do a chi-square. Or we could use the data frame for some other statistical analysis (e.g., logistic regressions, etc). Oh and for real we do not need to save all those intermediate objects.

2 Two subsetting operations

Since getting access to selected parts of matrices, data frames, etc, is so important, two quick exercises.

First, obtain the mean of all the genes for the third subject, but only for those genes with a p-value less than 0.01. Or, to put it another way, you want to see the average (mean) expression of the genes in subject 3, but only for the genes that have a p-value < 0.01 :

```
## [1] 0.01816454
```

And now the median of the second gene, but ONLY for males

```
## [1] -1.04533
```

3 Gene summaries by condition and sex

The next exercises all involve summaries of specific genes. Remember that genes were originally the rows of the data set. Since we will be using aggregate and similar, you will want to have genes as columns. How? Doing a transposition:

```
leuk.dat.t <- t(leuk.dat.m)
```

Make sure you understand what happened. For instance, do:

```
dim(leuk.dat.t)

## [1] 38 3051
```

and now check the first few columns (you do not want to use head, because it would show all 3051 columns)

```
leuk.dat.t[1:5, 1:6]

##           G1           G2           G3           G4           G5           G6
## V2 -1.45769 -0.75161  0.45695  3.13533  2.76569  2.64342
## V3 -1.39420 -1.26278 -0.09654  0.21415 -1.27045  1.01416
## V4 -1.42779 -0.09052  0.90325  2.08754  1.60433  1.70477
## V5 -1.40715 -0.99596 -0.07194  2.23467  1.53182  1.63845
## V6 -1.42668 -1.24245  0.03232  0.93811  1.63728 -0.36075
```

3.1 aggregate: The median of three genes by condition

Use aggregate to obtain the median of genes in positions 1, 2124, and 2600 (HK-1, PTEN, and the other gene from the figures in the exercises from last week)

```
##   type      G1      G2124      G2600
## 1  ALL -1.36149 -0.21672 -0.11015
## 2  AML -1.22961  1.57629  1.32525
```

3.2 aggregate: The median of three genes by condition and sex

Like the previous exercise, but you want to obtain those summaries by condition and sex:

```
##   type  sex      G1      G2124      G2600
## 1  ALL Female -1.368320  0.077270 -0.00771
## 2  AML Female -1.320115  1.718425  1.13218
## 3  ALL  Male -1.321430 -0.460435 -0.16511
## 4  AML  Male -1.089020  1.457100  1.55580
```

3.3 aggregate: The median of all the genes by condition and sex

This is like the previous one, but for **every** gene. So you do not want to print it. Assign the result of aggregate to an object, let's call it "all.median".

Now, show all the rows and the first 10 columns

```
##      type      sex      G1      G2      G3      G4      G5
## 1  ALL Female -1.368320 -1.032090 0.233810 0.641800 0.22853
## 2  AML Female -1.320115 -0.922255 0.147070 0.360200 -0.57654
## 3  ALL  Male -1.321430 -1.034565 0.266325 2.206875 1.71448
## 4  AML  Male -1.089020 -1.089020 0.046090 1.302970 1.01596
##      G6      G7      G8
## 1  0.968700 3.287650 3.039720
## 2 -0.164105 3.126385 2.901165
## 3  1.614675 3.087135 2.984870
## 4  0.512660 3.279340 2.961040
```

and the dimensions of "all.median"

```
dim(all.median)

## [1]      4 3053
```

3.4 aggregate: The mean and standard deviations of three genes by condition and sex

This is like the above ones, but we do not want just one statistic (the median); here we want the mean and standard deviation for all our three genes (1, 2124, 2600). And, if possible, with clear column names, so you will have to be explicit, when you pass a function to `aggregate`, about the names of the components you return. (In other words, DO NOT assign the return value of `aggregate` and change the column names; you have to play around with the names in the return object from the function you pass to `aggregate`)

```
##      type      sex      G1.mean      G1.sd G2124.mean      G2124.sd
## 1  ALL Female -1.3180008    0.1566979 -0.1633385    0.5815847
## 2  AML Female -1.1569633    0.3503542    1.6656317    0.3161407
## 3  ALL  Male -1.2283407    0.4253546 -0.4168571    0.8363303
## 4  AML  Male -0.3259880    1.1929901    1.4919120    0.4385808
##      G2600.mean      G2600.sd
## 1 -0.23134692    0.63891862
## 2  1.30763500    0.50586581
## 3 -0.09672786    0.78109489
## 4  1.42469200    0.40976226
```

3.5 by and aggregate: those three genes again, but now use "summary"

As it says. First, do not use `aggregate` now, but use `by`, and the function you want to have applied to each subset is `summary`, which will give you different summary statistics.

```
## type: ALL
## sex: Female
##      G1      G2124      G2600
## Min.   :-1.608   Min.    :-1.36832   Min.    :-1.40095
## 1st Qu.: -1.401   1st Qu.: -0.40103   1st Qu.: -0.66446
## Median :-1.368   Median : 0.07727   Median : -0.00771
## Mean   :-1.318   Mean    :-0.16334   Mean    :-0.23135
## 3rd Qu.: -1.216   3rd Qu.: 0.25025   3rd Qu.: 0.17697
## Max.   :-1.032   Max.     : 0.62632   Max.     : 0.57343
## -----
## type: AML
## sex: Female
```

```
##           G1           G2124           G2600
## Min.      :-1.4443    Min.      :1.228    Min.      :0.8156
## 1st Qu.   :-1.4104    1st Qu.   :1.463    1st Qu.   :0.9295
## Median    :-1.3201    Median    :1.718    Median    :1.1322
## Mean      :-1.1570    Mean      :1.666    Mean      :1.3076
## 3rd Qu.   :-0.8941    3rd Qu.   :1.810    3rd Qu.   :1.7284
## Max.      :-0.6646    Max.      :2.110    Max.      :1.9767
## -----
## type: ALL
## sex: Male
##           G1           G2124           G2600
## Min.      :-1.5678    Min.      :-1.4765    Min.      :-1.23051
## 1st Qu.   :-1.4275    1st Qu.   :-1.2253    1st Qu.   :-0.58571
## Median    :-1.3214    Median    :-0.4604    Median    :-0.16511
## Mean      :-1.2283    Mean      :-0.4169    Mean      :-0.09673
## 3rd Qu.   :-1.2341    3rd Qu.   : 0.3919    3rd Qu.   : 0.39772
## Max.      : 0.1763    Max.      : 0.5792    Max.      : 1.49820
## -----
## type: AML
## sex: Male
##           G1           G2124           G2600
## Min.      :-1.262    Min.      :1.054    Min.      :0.8136
## 1st Qu.   :-1.230    1st Qu.   :1.189    1st Qu.   :1.2176
## Median    :-1.089    Median    :1.457    Median    :1.5558
## Mean      :-0.326    Mean      :1.492    Mean      :1.4247
## 3rd Qu.   : 0.849    3rd Qu.   :1.576    3rd Qu.   :1.7369
## Max.      : 1.101    Max.      :2.183    Max.      :1.7997
```

(Of course, if you wanted, you could do that to all of the genes in the array)
Now, use aggregate:

```
##   type    sex  G1.Min. G1.1st Qu. G1.Median  G1.Mean G1.3rd Qu.
## 1  ALL Female -1.607670 -1.400950 -1.368320 -1.318001 -1.215830
## 2  AML Female -1.444340 -1.410423 -1.320115 -1.156963 -0.894055
## 3  ALL  Male -1.567830 -1.427512 -1.321430 -1.228341 -1.234088
## 4  AML  Male -1.261830 -1.229610 -1.089020 -0.325988  0.849050
##      G1.Max. G2124.Min. G2124.1st Qu. G2124.Median G2124.Mean
## 1 -1.032090 -1.3683200   -0.4010300   0.0772700 -0.1633385
## 2 -0.664650  1.2281400    1.4631700   1.7184250  1.6656317
## 3  0.176280 -1.4764900   -1.2252900  -0.4604350 -0.4168571
## 4  1.101470  1.0537100    1.1894700   1.4571000  1.4919120
##      G2124.3rd Qu. G2124.Max.  G2600.Min. G2600.1st Qu. G2600.Median
## 1      0.2502500  0.6263200 -1.40095000  -0.66446000  -0.00771000
## 2      1.8098725  2.1103600  0.81563000   0.92949500   1.13218000
## 3      0.3919050  0.5792000 -1.23051000  -0.58571500  -0.16511000
## 4      1.5762900  2.1829900  0.81357000   1.21756000   1.55580000
##      G2600.Mean G2600.3rd Qu.  G2600.Max.
## 1 -0.23134692   0.17697000  0.57343000
## 2  1.30763500   1.72844250  1.97669000
## 3 -0.09672786   0.39772000  1.49820000
## 4  1.42469200   1.73686000  1.79967000
```

Please, notice the differences in output. They are different kinds of objects, something you can see if you assign the output to an object and do, for instance

```
class(objeto)
```